

REPORT: CROSS-MODAL IMAGE MATCHING & POTENTIAL IMPROVEMENTS

Project: Siamese Network for Sentinel-1 (SAR) & Sentinel-2 (Optical) Imagery Alignment
Prepared by: Artem
Status: Baseline Achieved (92% Similarity Alignment)

1. Executive Summary

The implementation of the Siamese Neural Network using a ResNet-18 backbone and Triplet Margin Loss has successfully achieved its primary objective: bridging the domain gap between Synthetic Aperture Radar (SAR) and Optical RGB spectrums. The baseline model achieved a cross-modal verification score of **0.9208**, demonstrating that deep metric learning can effectively align structural land topologies despite severe radiometric noise, sensor differences, and speckle artifacts inherent in radar data. However, to transition this prototype into a production-grade geospatial application, several architectural and pipeline enhancements can be introduced.

2. Identified Bottlenecks in the Baseline

- **Global vs. Local Feature Mismatch:** The Siamese network optimizes a global image embedding. While SIFT and RANSAC successfully find local geometric matches in post-processing, the neural network itself does not explicitly learn local keypoint descriptors across modalities.
- **Radarsat Speckle Noise:** Sentinel-1 SAR imagery is highly susceptible to multiplicative speckle noise, which can distort fine-grained edge detection during the initial epochs.
- **Basic Negative Sampling:** Random negative sampling in the triplet generator can lead to 'easy triplets' after a few epochs, causing the loss gradients to flatten quickly or stall out.

3. Mathematical Approach & Architecture Updates

To suppress embedding collapse, output vectors are forced onto a unit hypersphere via L2 Normalization:

$$e = f(x) / ||f(x)||_2$$

The network currently minimizes the Triplet Margin Loss targets:

$$L(a, p, n) = \max(d(e_a, e_p) - d(e_a, e_n) + \text{margin}, 0)$$

4. Proposed Experimentation Roadmap

Improvement Type	Difficulty	Expected Impact	Resource Cost
Online Hard Negative Mining	Medium	High (Faster convergence)	Low
Pseudo-Siamese Architecture	Medium	Very High (Higher accuracy)	Medium (2x VRAM)
Lee/Frost SAR Despeckling	Low	Moderate (Cleaner features)	Low
LoFTR / SuperGlue Integration	High	Critical (Flawless local matching)	High (GPU Required)

5. Conclusion

The baseline project successfully proved the mathematical viability of cross-modal Siamese registration. Implementing Online Hard Negative Mining and migrating to a Pseudo-Siamese network architecture represent the highest-ROI next steps, promising to lift verification confidence, accelerate training stability, and deliver bulletproof localization accuracy.